Service-oriented middleware for mobile sensing

110015230, 110006876

Goals

Modern mobile devices are equipped with various complex sensors – accelerometers, location, video, audio, temperature, altitude, etc. Not all of them provide scientific-grade precision or update frequency. However these disadvantage are often compensated by ubiquity and availability of such devices, allowing unprecedented deployment scale for scientific experiments. We aim to create a piece of message-oriented middleware which allows data to be gained from mobile devices without the data recipient having to build a server architecture.

Design

Key protocols

As described in the brief, the sensor data producer is deployed on a mobile device. This dictated a number of constraints that had to be considered while developing the system; First, the technologies used to develop the producer had to be easy to deploy on a mobile device. In recent years, it has become increasingly popular to develop mobile web apps – applications that use the device's web browser as the execution environment. Such applications have an advantage of being easy to distribute and deploy on various platforms. HTML5 and Javascript, used to develop web apps, provide rich APIs for accessing sensor data – accelerometers, audio, video, etc. Thus, the project aimed to utilise technologies that are supported by browsers of modern mobile platforms.

WebSockets and WAMP

WebSockets is a protocol for full-duplex (two-way) communication through TCP, meaning various aspects of reliability and robustness are taken care of by the transport protocol (e.g. in-order guaranteed messages). Additionally, compared to regular HTTP requests, WebSockets have a low header overhead (2 bytes compared to kilobytes in HTTP). Two-way communication and low overhead (thus, decreased latency and bandwidth usage) are advantageous for a data streaming platform.

WebSockets are supported by most modern desktop and mobile browsers (through Javascript APIs, without needing any plugins), which makes it easy to use WebSockets-based communication in web-apps. Various libraries implement the WebSockets for other languages – Python, Java, Go, etc.

WAMP is a standard that builds upon WebSockets, providing some Pub-Sub and RPC functionality. There are currently two versions of WAMP – V1 and V2. The more stable V1 was selected due to the existence of many libraries implementing the standard.

[http://wamp.ws/implementations/wamp1/]

REST

Representational state transfer (REST) is a set of principles that uses Web standards (such as HTTP and URL) to structure application-to-application communication. The system operates through a REST API, the registry offering various methods to clients (e.g. '/register_producer', '/request_brokers') which are called via GET requests. This means that any language with the ability to send HTTP requests can use the system.

Components

Registry

The system has a single central registry which keeps track of which producers are available and where their broker is. The registry saves the current list of producers to disk whenever it is updated, allowing it to be rebuilt after a crash. Additionally, it can kill the brokers of producers which have disappeared for an extended period, and give new brokers to producers who request one due to their original broker disconnecting unexpectedly.

API calls

Register producer:

Sample request

```
GET <REGISTRY_URI>/register_producer?sensors=camera,microphone,location&lat=55.755826&lng=37.6173
```

Sample response {"broker_address": "ws://localhost:61227", "lat": "51.755826", "lng": "37.6173", "location": "Kursk-Voronezh - Kshenskiy avtodoroga, Kursk Oblast, Russia"}

List brokers GET /list_brokers

Broker

The broker is a simple server which subscribes to a single producer and allows multiple consumers to access this data by publishing it on different channels for each sensor, essentially acting as a proxy.

Producer

It was decided that a producer must send the data only to one source. Multiple consumers might be interested in the data provided by the producer, which would cause scalability issues in the case of direct producer-to-consumer communication. Implementing data broadcasting through the producer would not be acceptable, since mobile devices are often limited in bandwidth. The producer sends data to a single broker, which can then pass it on to many consumers. The producer is also able to recover if its broker dies, instantly reconnecting with the registry so it can commence publishing on a new broker.

Consumer

The consumer has a list of brokers (obtained from the registry via the REST API) to which it subscribes to get data from producers. This data is then written to file, although any kind of processing could be carried out.

Implementation

Components

Registry

Registry is used to: register and keep track of data providers, as well as providing addresses of relevant brokers to the inquiring consumers. A REST API is used for communicating with the registry. The popular Flask web framework is used to serve HTTP requests.

The registry maintains a list of brokers (and, thus, producers associated with them). For each producer-broker pair, we keep track of address of the broker, available sensors, current location, the time of last heartbeat and other information. At every change, the list is serialised and saved in a file (using Pickle). This allow to restore the state of the registry after it has crashed (on start, the contents of broker list are loaded from the serialised file).

The registry is responsible for spawning brokers for newly registered producers. Python's subprocess module was used to execute processes that would run brokers. The aim was to make the brokers independent of the registry. Generally, child processes are directly reliant upon the parent process. Thus, if the registry (parent) crashes, the brokers would be closed (children). Such behavior was suppressed by providing a number of special flags and commands:

- nohup – is a POSIX command that tells the process to ignore the Hangup signal
- preexec_fn=os.setpgrp – creates new process group and places the child process into it
- stdout=open('/dev/null', 'w'), stderr=open('logfile.log', 'a') – detaches the process from the console of the parent process
- close_fds = True – allows to prevent the child processes inheriting the parent's open ports (through file descriptors). This is important because otherwise the registry would not be able to access the default port after the restart, since the children are blocking it.

As a result, the brokers are running completely independently of the producer process. This allows the to restart the registry without disrupting the brokers (and vice versa).

Broker

The broker is spawned by the registry. A list of sensors available to the producer is passed to the broker where it, following the WAMP standard, is registered as a channel for pub-sub. The broker then proceeds to fulfill its main duty – using WebSockets, listen to incoming messages from producer and broadcast them to the subscribed consumers.

It should be noted that because WebSockets allow bi-direction communication, there is very little difference for the broker between consumer and producer. In future version, this feature

Producer

Consumer

Implementation justification

No-DB

Initially, we have considered using a database as an intermediate storage for messages coming from the producer to the consumer. Databases are often considered to be robust, scalable and potentially distributed. However, after considering the pros and cons of such approach, we decided against it.

Using a single database for multiple producers introduces an obvious single point of failure. Distributing databases through sharding or other techniques can potentially solve this problem, however it is not not trivial – corporations employ software reliability engineers, buy expensive hardware, work on the underlying DB management system to achieve this task. Various providers offer hosted DB installations with load balancing, uptime guarantees and considerable cost of service (Google Cloud SQL). However, we believe that relying too much on a third party service would defeat the purpose of the practical.

Thus, we decided to implement our pub-sub architecture without relying on a database management system.

Load balancing

As previously mentioned, a separate broker is spawned per each producer. This allows to avoid maintaining multiple producer-to-broker connections, thus maintaining a constant CPU and network load for the producer (important, since it is a mobile device with limited connectivity and computational power).

This approach works particularly well, if the interest for the sensors is homogeneously distributed. However, some producers might be more interesting to consumers than the others.

There are several techniques that would allow to further improve load balancing of the brokers.

- Chaining brokers –
- Hardware based —

Testing

Broker scalability

Testing was conducted to check the scalability of the broker processes. An experiment was set up: 1. The registry is started up 2. A single producer is started 3. Multiple consumers connect and request the data from that producer.

Since the producer-to-consumer communication is done through a single broker, we are able to tests how well does the broker cope with with serving multiple consumers. All components of the system (registry, producer, broker, consumers) are running on the same computer, in order to avoid measuring time used for the network communication, rather than the time used for processing.

The turnaround time between the message sent by the producer and it being received by the consumer is measured (in seconds).

Robustness

Temporary Interruption

As the system is built on top of TCP, messages are guaranteed to arrive in order by default. Examples of this can be seen in the file shown when testing the rebuilding of the registry.

Removing Dead Producers from the Registry

If a producer unexpectedly dies (or loses connectivity, e.g. because the device went into a tunnel) the registry will stop receiving its heartbeat. After a timeout period the producer will be considered gone for good, meaning its broker must be killed and its details removed from the registry. To test this, the registry and three producers (in Moscow) were run. After this stage, the list of brokers looked as follows:

Figure 1: Initial list of producers

Two of the producers were then killed (by sending SIGINT from the terminal). After this, the list of producers was still the same, as the registry only cleans up the producer list either when a consumer requests data from new producers or when it is explicitly told to (via the /clean_producers API call). After this call was made, the list looked as follows:

When the command is called, the registry prints the ID of any producer which is being removed to its output stream, which is shown below just for more confirmation that it actually did happen.

Reconnecting With Producers Whose Broker has Died

**Rebuilding the Registry After a Crash**

Reproducing tests

Future work

Any system can always be more scalable and robust, and thus there is much that could be done to improve the system.

There are a few ways the current implementation of the system can be further improved:

The registry, based on the popular Flask framework, currently can withstand many hundreds of requests at once. However, at the moment it is running on the development server (server software, not hardware) which is not optimised for concurrent requests. Deploying the registry on a combination of Nginx (high-performance HTTP server) and Gunicorn (WSGI HTTP server) would allow to "hide" multiple registry processes, responding to the same port. This would considerably improve the scalability of the registry for concurrent requests. Unfortunately, deploying Nginx and Gunicorn was problematic on the school's servers. For now, Flask's development server was used.

File Edit View History Bookmarks Tools Help

□ http://127.0.0....00/list_brokers ✕ | □ http://127.0.0...lean_producers ✕ | ✛

127.0.0.1:5000/list_brokers   ✕ ∨ ⟳ 8 ∨ Google ⟳ ⌂ ⚟ ∨

1

['MGAQeewNZ97NK65Oj0GtQnv6c1FSHK', ['camera', 'microphone', 'location'],
'55.76', '37.62', datetime.datetime(2014, 4, 2, 22, 3, 8, 547511),
'ws://127.0.0.1:57792', 17102]

Figure 2: Cleaned list of producers

Removing producer 0HCnJOkQDjUG0ZAdBrbXyQyxiQfVUL
Removing producer Q83u7A9AkyS1RXCWD0NAGy84hAhHAi
127.0.0.1 - - [02/Apr/2014 22:03:12] "GET /clean_producers HTTP/1.1" 200 -

Figure 3: Producers which are being removed

With more time, we would have liked to implement distributed brokers (i.e. running them on other machines), with the possibility for brokers to subscribe to and publish data from multiple producers to allow the system to operate with fewer processes.

Conclusion

A robust message-oriented middleware for mobile sensing was developed. The system is based on the pub-sub paradigm. It consists of several components – producers, consumers, registry, brokers. Components communicate using modern standards: WebSockets + WAMP and REST. According to the component specification, a Python implementation was created, although other languages could easily be used to implement producers or consumers. The implementation was testing, showcasing good scalability, robustness, and ability to recover from crashes.